

HARVARD EXTENSION SCHOOL

EXT CSCI E-106 Model Data Class Group Project Template

Author: Dinesh Bedathuru Author: Brian Calderon
Author: Jeisson Hernandez Author: Hao Fu Author: Derek Rush
Author: Jeremy Tajonera Author: Catherine Tully

08 May 2025

In this project, our aim is to classify the probability of a passenger surviving the Titanic crash of 1912. We used a variety of linear and non-linear models to deduce the most accurate model and provide long-term stability in our predictions.

Table of contents

1	Instructions:	2
2	Executive Summary	3
3	Introduction	3
4	Description of the data and quality	4
4.1	Loading the data	4
4.2	Removing un-needed columns	4
4.3	Data Augmentation	4
4.4	Initial Check for Missing values	5
4.5	Imputing data	5
4.6	Dummifying Columns:	7
4.7	Remove NA rows and deck_G	8
4.8	Divide into Test / Train	8
4.9	EDA	9
5	Model Development Process	14
6	Model Performance Testing	17
7	Challenger Models	26
8	Model Limitation and Assumptions	30
9	Ongoing Model Monitoring Plan	32
10	Conclusion	33
	Appendix A: Check if ‘sibsp’ and ‘parch’ should be continuous or categorical	33
	References	47

List of Figures

1	Percentage of Missing Values	5
2	Percentage of Missing Values after Imputation	7
3	Histogram of Deck_G	8
4	Histogram of survival vs age	9
5	Barchart of survival	10
6	Barchart of survival vs Num. of siblings	10
7	Scatter plots of all variables in train data	12
8	Correlation Matrix	14
9	4x4 standard plots for log. model	17

List of Tables

2	Summary Stats. of base log. model	16
3	Comparison of all log models performance	26
6	Summary of Diagnostic tests on Filtered vs Base models	31
7	VIF Results for Base Logistic Regression Model	31
8	VIF Results for Filtered Logistic Regression Model	31
9	Comparison of stats between base and filtered log models:	32

Classify whether a passenger on board the maiden voyage of the RMS Titanic in 1912 survived given their age, sex and class. Sample-Data-Titanic-Survival.csv to be used in the Final Project

Variable	Description
pclass	Passenger Class (1 = 1st; 2 = 2nd; 3 = 3rd)
survived	Survival (0 = No; 1 = Yes)
name	Name
sex	Sex
age	Age
sibsp	# of siblings / spouses aboard the Titanic
parch	# of parents / children aboard the Titanic
ticket	Ticket number
fare	Passenger fare
cabin	Cabin number
embarked	Port of Embarkation (C = Cherbourg; Q = Queenstown; S = Southampton)
boat	Lifeboat ID, if passenger survived
body	Body number (if passenger did not survive and body was recovered)
home.dest	The intended home destination of the passenger

1 Instructions:

0. Join a team with your fellow students with appropriate size (Up to Nine Students total) If you have not group by the end of the week of April 11 you may present the project by yourself or I will randomly assign other stranded student to your group. I will let know the final groups in April 11.

1. Load and Review the dataset named “Titanic_Survival_Data.csv” 2. Create the train data set which contains 70% of the data and use set.seed (15). The remaining 30% will be your test data set.

3. Investigate the data and combine the level of categorical variables if needed and drop variables as needed. For example, you can drop id, Latitude, Longitude, etc.
4. Build appropriate model to predict the probability of survival.
5. Create scatter plots and a correlation matrix for the train data set. Interpret the possible relationship between the response.
6. Build the best models by using the appropriate selection method. Compare the performance of the best logistic linear models.
7. Make sure that model assumption(s) are checked for the final model. Apply remedy measures (transformation, etc.) that helps satisfy the assumptions.
8. Investigate unequal variances and multicollinearity.
9. Build an alternative to your model based on one of the following approaches as applicable to predict the probability of survival: logistic regression, classification Tree, NN, or SVM. Check the applicable model assumptions. Explore using a negative binomial regression and a Poisson regression.
10. Use the test data set to assess the model performances from above.
11. Based on the performances on both train and test data sets, determine your primary (champion) model and the other model which would be your benchmark model.
12. Create a model development document that describes the model following this template, input the name of the authors, Harvard IDs, the name of the Group, all of your code and calculations, etc..

Due Date: May 12 2025 1159 pm hours EST Notes No typographical errors, grammar mistakes, or misspelled words, use English language All tables need to be numbered and describe their content in the body of the document All figures/graphs need to be numbered and describe their content All results must be accurate and clearly explained for a casual reviewer to fully understand their purpose and impact Submit both the RMD markdown file and PDF with the sections with appropriate explanations. A more formal.

2 Executive Summary

This section will describe the model usage, your conclusions and any regulatory and internal requirements. In a real world scenario, this section is for senior management who do not need to know the details. They need to know high level (the purpose of the model, limitations of the model and any issues).

3 Introduction

This section needs to introduce the reader to the problem to be resolved, the purpose, and the scope of the statistical testing applied. What you are doing with your prediction? What is the purpose of the model? What methods were trained on the data, how large is the test sample, and how did you build the model?

The Titanic was a British-registered ship that set sail on its maiden voyage on April 10th, 1912 with 2,240 passengers and crew on board. On April 15th, 1912, the ship struck an iceberg, split in half, and sank to the bottom of the ocean (National Oceanic and Atmospheric Administration (NOAA), 2023). In this report, we are going to analyze the data in the Titanic.csv file and use it to determine the best model for predicting whether someone on board would live or die. By creating this model, we hope to understand what factors a passenger could have taken into account in order to reduce their risk of death during the trip. We cleaned the data and split into into a train/test split in order to properly train our models. We created simple linear models, multivariate linear models, logistic models (both binomial and poisson), a regression tree, and a neural network model. The train sample size was 916 data points (70.03%) and the test sample size was 392 data points (29.97%). We built the

models after examining the data and determining which predictor variables we thought would be most relevant for survival rate. Once we had our variables and training data, we created the models and examined the performance of the models on both training and testing data to determine if they were robust. We also examined if the model assumptions appeared to hold for each model.

4 Description of the data and quality

Based on the data cleaning we were able to only remove 2 rows from the data set. We used median imputation as well as KNN for various columns. We also dummified several categorical columns. We found that leaving sibsp and parch as continuous as opposed to categorical increased their contributions to the model performance [Appendix A](#). Further, we also extracted the deck number and found that removing deck_G from the model increased its performance.

4.1 Loading the data

```
odata <- read.csv("../data/Titanic_Survival_Data.csv")
cat("Size of entire data set:", nrow(odata), "\n")
```

Size of entire data set: 1310

4.2 Removing un-needed columns

Name: Removing because names have no inference on survival (inference)

ticket: Ticket No. will also likely not have an influence in survival

boat: This is highly correlated to the survival dependant variable since people who made it on a boat likely survived

body: This is highly correlated to the survival dependant variable since people who's body was recovered did not survive.

home.dest: The destination likely has nothing to do with the survival

```
data.clean = odata[, !(names(odata) %in% c("name", "ticket", "boat", "body", "home.dest"))]
```

4.3 Data Augmentation

We extracted the deck letter from the cabin since it could potentially correlate to the survival.

```
#Extract deck letter from cabin
data.clean$deck <- substr(data.clean$cabin, 1,1)
# Remove cabin col:
data.clean$cabin <- NULL
```

4.4 Initial Check for Missing values

We see that age and deck have the most amount of missing data, therefore we proceed to impute them.

```
print(plot_missing_barchart(data.clean))
```

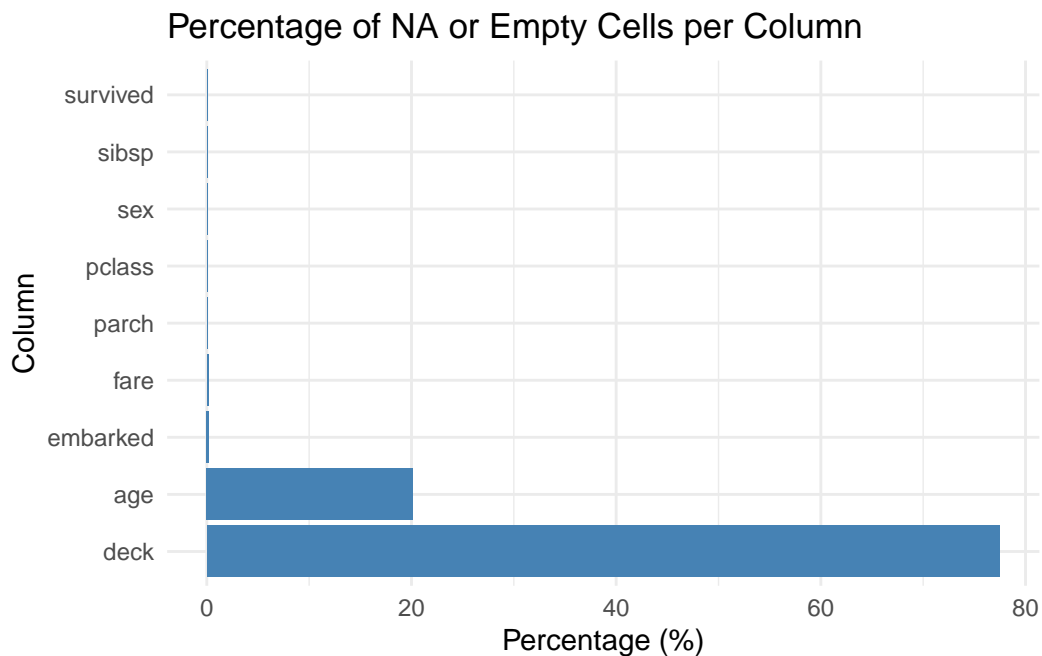


Figure 1: Percentage of Missing Values

4.5 Imputing data

Below we impute Age using the median value in that column.

For deck we use KNN to impute the missing deck values.

After imputing these two columns we can see that the largest amount of missing data is ~0.2% which is quite small and can be removed.

```
# ---- Age----
#Replace NAs in age column with Median value
median_age <- median(data.clean$age, na.rm = TRUE)
data.clean <- data.clean %>%
  mutate(age = ifelse(is.na(age), median_age, age))

# ---- deck----
# For deck, since its a category, we decided to use KNN to impute the column:

# Install if not already installed
# install.packages("VIM")
library(VIM)
```

Warning: package 'VIM' was built under R version 4.4.3

Loading required package: colorspace

Loading required package: grid

VIM is ready to use.

Suggestions and bug-reports can be submitted at: <https://github.com/statistikat/VIM/issues>

Attaching package: 'VIM'

The following object is masked from 'package:datasets':

sleep

```
# Replace "" with NA in the 'deck' column
data.clean$deck[data.clean$deck == ""] <- NA

# Convert 'cabin' to factor
data.clean$deck <- as.factor(data.clean$deck)

# Apply kNN imputation just to Cabin column
data.clean <- kNN(data.clean, variable = "deck", k = 5)

# Check that NAs were imputed
# sum(is.na(data.clean$deck))          # Original
# sum(is.na(data.clean$imputed$deck)) # After

# Remove indicator col:
data.clean$deck_imp <- NULL
```

```
#####
#           Check for Missing values after Imputation           #
#####

plot_missing_barchart(data.clean)
```




Figure 2: Percentage of Missing Values after Imputation

4.6 Dummifying Columns:

We dummify pclass, sex, embarked and deck. We leave sibsp and parch as continuous variables as we observed that dummifying these columns leads to smaller significance [Appendix A](#), whilst leaving them as continuous maximizes their contributions to the models explanatory power.

```
# Dummifying pclass:
data.clean$pclass_1 = ifelse(data.clean$pclass == 1, 1, 0)
data.clean$pclass_2 = ifelse(data.clean$pclass == 2, 1, 0)

# Dummifying sex:
data.clean$sex_M = ifelse(data.clean$sex == 'male', 1, 0)

# Dummifying embarked:
data.clean$embarked_C = ifelse(data.clean$embarked == 'C', 1, 0)
data.clean$embarked_Q = ifelse(data.clean$embarked == 'Q', 1, 0)

# Dummifying deck:
data.clean$deck_A = ifelse(data.clean$deck == 'A', 1, 0)
data.clean$deck_B = ifelse(data.clean$deck == 'B', 1, 0)
data.clean$deck_C = ifelse(data.clean$deck == 'C', 1, 0)
data.clean$deck_D = ifelse(data.clean$deck == 'D', 1, 0)
data.clean$deck_E = ifelse(data.clean$deck == 'E', 1, 0)
data.clean$deck_F = ifelse(data.clean$deck == 'F', 1, 0)
data.clean$deck_G = ifelse(data.clean$deck == 'G', 1, 0)

# Removing Dummified cols:
data.clean = subset(data.clean, select = -c(pclass, sex, embarked, deck))
```

4.7 Remove NA rows and deck_G

Below we remove NA rows, which turned out to be only 2 after proper cleaning and imputation. We also removed deck_G as we observed that it has a large skew in the data distribution with only 13 people allocated in this deck. It was observed that this variable lead to erroneous predictions in the model.

```
# Plot histogram of the 'values' column
hist(data.clean$deck_G,
     main = "Histogram of Values",
     xlab = "Values",
     col = "skyblue",
     border = "white")
```




Figure 3: Histogram of Deck_G

```
# Removing deck_G col:
data.clean = subset(data.clean, select = -c(deck_G))

data.clean = na.omit(data.clean)
cat(nrow(odata) - nrow(data.clean), 'rows were removed from original dataset')
```

2 rows were removed from original dataset

4.8 Divide into Test / Train

Finally we divide into 70% training data and 30% test data.

```
set.seed(1023)
train_indices = sample(1 : nrow(data.clean), size = 0.7005*nrow(data.clean), replace = FALSE)
train = data.clean[train_indices,]
test = data.clean[-train_indices,]
cat("We are using:", nrow(train)/nrow(data.clean) * 100, '% of the data for training')
```


We are using: 70.03058 % of the data for training

4.9 EDA

Using the training data set we use a variety of method to draw some initial conclusions:

- Histogram: Showing that more people in their late teens up to late thirties survived.
- Bar chart showing that more people died than survived
- Bar chart showing that a higher number of people survived when they had less siblings on board.
- Correlation matrix shows that sex and Deck_F are highly negatively correlated to survival. There is a soft positive correlation to pclass_1.
- There is a high correlation between pclass_1 and fare, this justifies that one of these predictors can potentially be removed.
- The scatter plots did not give us much more information on the relation between the predictors and the dependent variable.

```
# Histogram showing that more people in their late teens up to late thirties survived.  
ggplot(train, aes(age)) +  
  geom_histogram(bins=30)
```




Figure 4: Histogram of survival vs age

```
# Bar chart showing that more people died than survived  
ggplot(train, aes(survived)) +  
  geom_bar()
```




Figure 5: Barchart of survival

```
# Bar chart showing that a higher number of people survived when they had less
# siblings on board.
ggplot(train, aes(sibsp, survived)) +
  geom_bar(stat='identity')
```




Figure 6: Barchart of survival vs Num. of siblings

```
cor(train)
```

survived

age

sibsp

parch

fare

survived	1.00000000	-0.04599216	-0.027543645	0.095720424	0.249620417
age	-0.04599216	1.00000000	-0.154952394	-0.122885529	0.164477215
sibsp	-0.02754365	-0.15495239	1.00000000	0.355216328	0.165296185
parch	0.09572042	-0.12288553	0.355216328	1.00000000	0.209878267
fare	0.24962042	0.16447722	0.165296185	0.209878267	1.00000000
pclass_1	0.28432063	0.34847627	-0.022306615	-0.016451014	0.591163043
pclass_2	0.06118909	0.01898825	-0.069089243	-0.017696720	-0.125637353
sex_M	-0.53377358	0.05385130	-0.127066747	-0.243501842	-0.201776401
embarked_C	0.14757622	0.05701305	-0.061215242	-0.001165488	0.272087792
embarked_Q	-0.02542950	-0.03373092	-0.060221066	-0.093253422	-0.126992379
deck_A	0.02818626	0.12375210	-0.063208812	-0.046447535	0.060156186
deck_B	0.17102793	0.11384754	-0.018056185	0.073894445	0.458426488
deck_C	0.15622066	0.15548653	0.032030626	-0.054359774	0.314379308
deck_D	0.11169173	0.07390772	-0.003562241	0.002702881	0.007239167
deck_E	0.32472414	0.11407881	-0.082487921	-0.030305432	-0.017794021
deck_F	-0.47931430	-0.28313314	0.052474332	0.010856728	-0.420674788
	pclass_1	pclass_2	sex_M	embarked_C	embarked_Q
survived	0.28432063	0.06118909	-0.53377358	0.147576220	-0.02542950
age	0.34847627	0.01898825	0.05385130	0.057013047	-0.03373092
sibsp	-0.02230662	-0.06908924	-0.12706675	-0.061215242	-0.06022107
parch	-0.01645101	-0.01769672	-0.24350184	-0.001165488	-0.09325342
fare	0.59116304	-0.12563735	-0.20177640	0.272087792	-0.12699238
pclass_1	1.00000000	-0.30255006	-0.11706927	0.285611052	-0.16373859
pclass_2	-0.30255006	1.00000000	-0.01057862	-0.143425262	-0.14105035
sex_M	-0.11706927	-0.01057862	1.00000000	-0.040151739	-0.09800681
embarked_C	0.28561105	-0.14342526	-0.04015174	1.00000000	-0.16362864
embarked_Q	-0.16373859	-0.14105035	-0.09800681	-0.163628636	1.00000000
deck_A	0.27817486	-0.07214490	0.05726877	0.179563214	-0.06045425
deck_B	0.42157158	-0.12290262	-0.10513911	0.197253957	-0.08030780
deck_C	0.49783147	-0.09377853	-0.08157273	0.166128834	-0.07869333
deck_D	0.11604898	0.04117623	-0.06599395	0.288111512	-0.09280535
deck_E	0.06047826	-0.03335428	-0.07748047	-0.085668330	-0.08697346
deck_F	-0.70753043	0.15770708	0.15681273	-0.353103247	0.22529873
	deck_A	deck_B	deck_C	deck_D	deck_E
survived	0.02818626	0.17102793	0.15622066	0.111691728	0.32472414
age	0.12375210	0.11384754	0.15548653	0.073907716	0.11407881
sibsp	-0.06320881	-0.01805619	0.03203063	-0.003562241	-0.08248792
parch	-0.04644753	0.07389444	-0.05435977	0.002702881	-0.03030543
fare	0.06015619	0.45842649	0.31437931	0.007239167	-0.01779402
pclass_1	0.27817486	0.42157158	0.49783147	0.116048985	0.06047826
pclass_2	-0.07214490	-0.12290262	-0.09377853	0.041176226	-0.03335428
sex_M	0.05726877	-0.10513911	-0.08157273	-0.065993951	-0.07748047
embarked_C	0.17956321	0.19725396	0.16612883	0.288111512	-0.08566833
embarked_Q	-0.06045425	-0.08030780	-0.07869333	-0.092805348	-0.08697346
deck_A	1.00000000	-0.04808710	-0.06241531	-0.055570446	-0.06660453
deck_B	-0.04808710	1.00000000	-0.08291288	-0.073820128	-0.08847788
deck_C	-0.06241531	-0.08291288	1.00000000	-0.095815832	-0.11484105
deck_D	-0.05557045	-0.07382013	-0.09581583	1.00000000	-0.10224685
deck_E	-0.06660453	-0.08847788	-0.11484105	-0.102246852	1.00000000
deck_F	-0.23590750	-0.31338100	-0.40675710	-0.362149521	-0.43405805
	deck_F				
survived	-0.47931430				
age	-0.28313314				

```

sibsp      0.05247433
parch      0.01085673
fare       -0.42067479
pclass_1   -0.70753043
pclass_2    0.15770708
sex_M       0.15681273
embarked_C -0.35310325
embarked_Q  0.22529873
deck_A     -0.23590750
deck_B     -0.31338100
deck_C     -0.40675710
deck_D     -0.36214952
deck_E     -0.43405805
deck_F      1.00000000

```

```
pairs(train[c(1:4,7,13)])
```

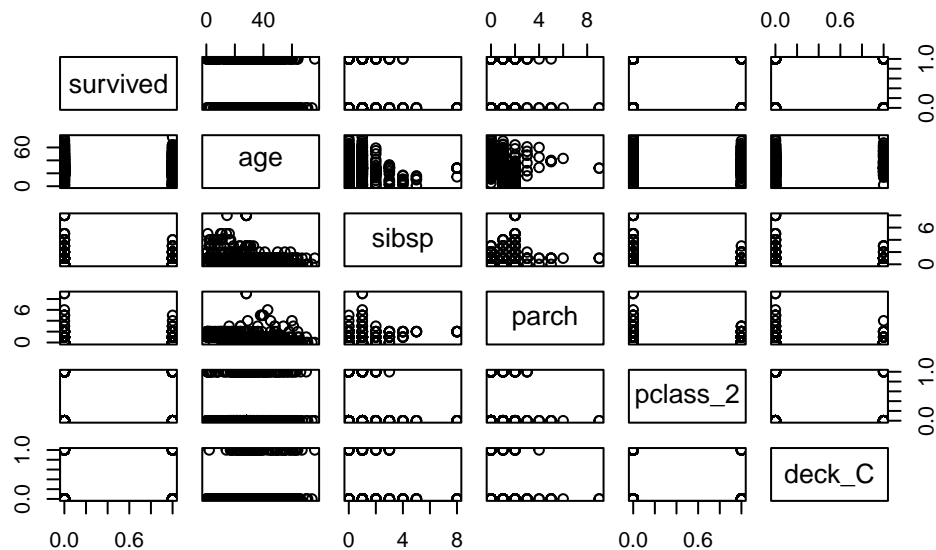


Figure 7: Scatter plots of all variables in train data

```
# Since this data is mainly categorical, the scatterplot and correlation matrix are not very useful.
```

(Statology, 2025) is used to develop the correlation values between our categorical columns. This describes the use of psych and rcompanion.

```

#install.packages("psych")
library(psych) # [Statology2025] to understand how this works

```

Warning: package 'psych' was built under R version 4.4.3

Attaching package: 'psych'

The following objects are masked from 'package:ggplot2':

%%, alpha

```
tetrachoric(train[, c("survived", "sex_M")])
```

Call: tetrachoric(x = train[, c("survived", "sex_M")])

tetrachoric correlation

	srvvd	sex_M
survived	1.00	
sex_M	-0.75	1.00

with tau of

	sex_M
survived	
0.26	-0.35

```
tetrachoric(train[, c("survived", "pclass_1")])
```

Call: tetrachoric(x = train[, c("survived", "pclass_1")])

tetrachoric correlation

	srvvd	pcl_1
survived	1.00	
pclass_1	0.46	1.00

with tau of

	pclass_1
survived	
0.26	0.69

```
tetrachoric(train[, c("survived", "pclass_2")])
```

Call: tetrachoric(x = train[, c("survived", "pclass_2")])

tetrachoric correlation

	srvvd	pcl_2
survived	1.00	
pclass_2	0.11	1.00

with tau of

	pclass_2
survived	
0.26	0.77

```
#install.packages("rcompanion")
```

```
library(rcompanion) # Reference 4 to understand how this works.
```

Warning: package 'rcompanion' was built under R version 4.4.3

Attaching package: 'rcompanion'

The following object is masked from 'package:psych':

phi

```
cramerV(train$survived, train$sex)
```

Cramer V
0.5338

```
library(corrplot)
```

corrplot 0.95 loaded

```
cor_matrix <- cor(train)#[,1]  
corrplot(cor_matrix, method = "circle")
```




Figure 8: Correlation Matrix

5 Model Development Process

The data was properly cleaned and divided into train/test in the prior section.

Here we train a binary model. The Q-Q plot shows that the residuals are indeed normally distributed so a transformation is potentially not necessary.

The statistical comparison between test and train data shows that the model is very stable with an accuracy of ~84% for both.

We also analyzed the VIF and we see that there is high degree of correlation between the decks, this provides justification to remove some of the decks as predictors.

```
library(car)
```

Warning: package 'car' was built under R version 4.4.3

Loading required package: carData

Warning: package 'carData' was built under R version 4.4.3

Attaching package: 'car'

The following object is masked from 'package:psych':

logit

The following object is masked from 'package:dplyr':

recode

The following object is masked from 'package:purrr':

some

```
# Log model on train data:
lmod <- glm(survived ~ ., family = binomial, data = train)
# summary(lmod)
vif(lmod)
```

age	sibsp	parch	fare	pclass_1	pclass_2	sex_M
1.493399	1.258340	1.281859	1.727034	4.267734	1.560100	1.604117
embarked_C	embarked_Q	deck_A	deck_B	deck_C	deck_D	deck_E
1.513511	1.397623	4.532575	4.848833	7.604802	6.662169	6.091048
deck_F						
14.919819						

```
y_hat_log_train <- predict(lmod, data = train, type="response")
predictions_log_train <- ifelse(y_hat_log_train > 0.5, 1, 0)
```

```
y_hat_log_test <- predict(lmod, newdata = test, type="response")
predictions_log_test <- ifelse(y_hat_log_test > 0.5, 1, 0)
```

```
confusion_matrix_log_train <- confusionMatrix(as.factor(predictions_log_train), as.factor(train$survived))
base.model.accuracy = confusion_matrix_log_train$overall['Accuracy']
base.model.f1 = confusion_matrix_log_train$byClass['F1']
base.model.train.summary = data.frame(
  Accuracy = base.model.accuracy,
  F1 = base.model.f1
)
row.names(base.model.train.summary) <- 'base.model.train'
```

```

confusion_matrix_log_test <- confusionMatrix(as.factor(predictions_log_test), as.factor(test$survived))
base.model.accuracy = confusion_matrix_log_test$overall['Accuracy']
base.model.f1 = confusion_matrix_log_test$byClass['F1']
base.model.test.summary = data.frame(
  Accuracy = base.model.accuracy,
  F1 = base.model.f1
)
row.names(base.model.test.summary) <- 'base.model.test'

data.frame(rbind(base.model.train.summary, base.model.test.summary))

```

Table 2: Summary Stats. of base log. model

	Accuracy	F1
base.model.train	0.8471616	0.8071625
base.model.test	0.8367347	0.7730496

```
summary(lmod)
```

Call:

```
glm(formula = survived ~ ., family = binomial, data = train)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)	
(Intercept)	5.3284111	0.8287318	6.430	1.28e-10	***
age	-0.0557904	0.0090406	-6.171	6.78e-10	***
sibsp	-0.2761926	0.1235912	-2.235	0.02544	*
parch	0.0250325	0.1186751	0.211	0.83294	
fare	0.0001907	0.0023082	0.083	0.93415	
pclass_1	0.8206018	0.4454290	1.842	0.06544	.
pclass_2	1.8580626	0.2991637	6.211	5.27e-10	***
sex_M	-3.2522858	0.2555063	-12.729	< 2e-16	***
embarked_C	0.6313152	0.2939442	2.148	0.03173	*
embarked_Q	0.8927717	0.3642358	2.451	0.01424	*
deck_A	-1.9122914	0.9044640	-2.114	0.03449	*
deck_B	-1.2596191	0.9199589	-1.369	0.17093	
deck_C	-1.6453843	0.8501358	-1.935	0.05294	.
deck_D	-2.3295283	0.8315684	-2.801	0.00509	**
deck_E	0.4168948	0.8021132	0.520	0.60324	
deck_F	-4.3018644	0.7733641	-5.563	2.66e-08	***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 1230.15 on 915 degrees of freedom

Residual deviance: 635.57 on 900 degrees of freedom

AIC: 667.57

Number of Fisher Scoring iterations: 6

```
par(mfrow=c(2,2))
plot(lmod)
```

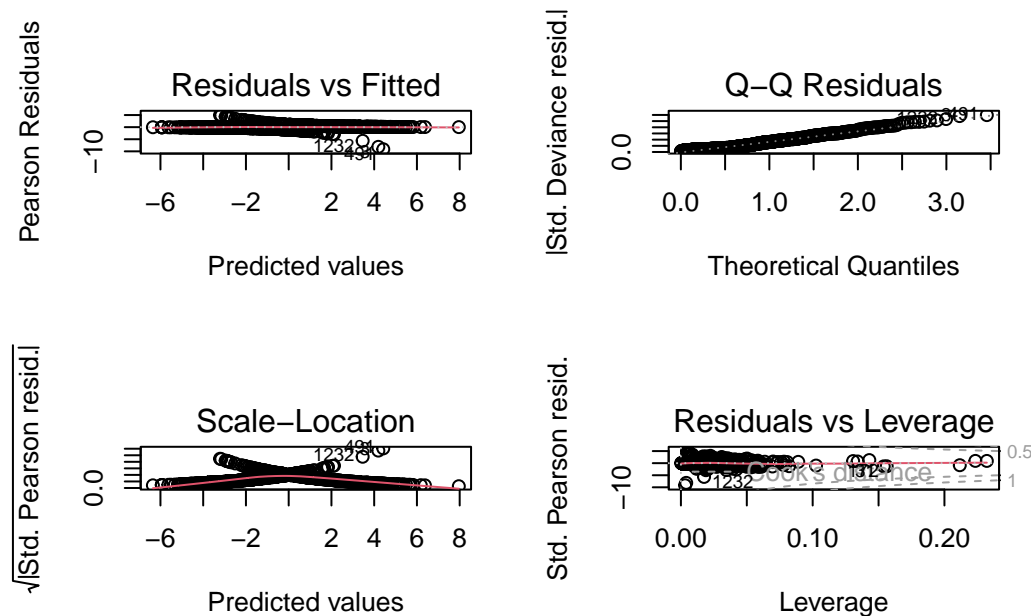


Figure 9: 4x4 standard plots for log. model

```
# Plots show that a linear model is not appropriate for this data.
```

6 Model Performance Testing

We compare four different models: 1) Base log. model 2) Model with insign. pred. removed 3) Stepwise model 4) Model with high vif pred. removed

When comparing the accuracy and F1 score on all models the base model was still the highest performer and we decided to use that as the champion model until now.

```
#####
#           Function to remove Insig. Predictors one by one           #
#####

backward_eliminate = function(model, alpha = 0.05) {
  repeat {
    d1 = drop1(model, test = "F")

    # Get p-values excluding intercept row
    pvals = d1$`Pr(>F)`[-1]
```

```

# Stop if all predictors are significant or only intercept left
if( all(is.na(pvals)) || max(pvals, na.rm = TRUE) <= alpha ){
  print("all variable are significant")
  break
}

# Remove the term with max p-value
term_to_remove = rownames(d1)[-1][which.max(pvals)]
cat("Removing:", term_to_remove, "with p-value", max(pvals, na.rm = TRUE), "\n")
model = update(model, paste(". ~ . -", term_to_remove))
}
return(model)
}

```

```

#####
#                               Function to remove Cooks Outliers                               #
#####
# model_formula: A formula object, e.g. PSA.level ~ .
# data: A data frame containing the variables in the model.
# threshold: A numeric value indicating the Cook's D threshold (default 0.5).
# print: If TRUE (default) will print the rows beign removed.
# returns: A list with the final model and the filtered dataset.
#
# Example usage:
# result = remove_cooks_outliers(PSA.level ~ ., mydata)
# summary(result$model)
# str(result$filtered_data)
remove_cooks_outliers = function(model_formula, data, threshold = 0.5,
                                  print = TRUE)
{

  all_high_cd_rows = data.frame() # to store all removed rows

  repeat {
    model = glm(model_formula, family = binomial, data = data)
    cooksD = cooks.distance(model)
    high_cd_indices = which(cooksD > threshold)

    if (length(high_cd_indices) == 0) { # If there are no more outliers
      break
    }

    if (print == TRUE){
      cat("Removing rows with Cook's D >", threshold, ":\n", high_cd_indices, "\n")
    }
    # Save these outliers before removing them
    high_cd_rows = data[high_cd_indices, ]
    all_high_cd_rows = rbind(all_high_cd_rows, high_cd_rows)

    # Update data by removing high CD rows for next iteration.
    data = data[-high_cd_indices, ]
  }
}

```

```

final_model = glm(model_formula, family = binomial, data = data)
return(list(model = final_model, filtered_data = data, high_cd_data = all_high_cd_rows))
}

```

```

#####
#                               Model with insig. variables removed                               #
#####
binary.model.filtered = backward_eliminate(lmod)

```

Warning in drop1.glm(model, test = "F"): F test assumes 'quasibinomial' family

Removing: fare with p-value 0.9215538

Warning in drop1.glm(model, test = "F"): F test assumes 'quasibinomial' family

Removing: parch with p-value 0.7803802

Warning in drop1.glm(model, test = "F"): F test assumes 'quasibinomial' family

Removing: deck_E with p-value 0.5506678

Warning in drop1.glm(model, test = "F"): F test assumes 'quasibinomial' family

[1] "all variable are signifcant"

```

#####
#                               step wise model with insig. variables removed                               #
#####
library(olsrr)

```

Warning: package 'olsrr' was built under R version 4.4.3

Attaching package: 'olsrr'

The following object is masked from 'package:MASS':

cement

The following object is masked from 'package:datasets':

rivers

```

ols_step_both_p(lmod,p_enter=0.1,p_remove=0.05,details=FALSE)

```

Stepwise Summary

Step	Variable	AIC	SBC	SBIC	R2	Adj. R2
0	Base Model	1293.366	1303.006	-28281799.310	0.00000	0.00000
1	sex_M (+)	988.183	1002.643	-55309468.385	0.28491	0.28413
2	deck_F (+)	757.429	776.709	-91941771.510	0.44537	0.44415
3	age (+)	725.480	749.580	-99004330.024	0.46555	0.46379
4	pclass_2 (+)	696.998	725.918	-105806102.781	0.48304	0.48077
5	deck_E (+)	673.931	707.671	-111746149.031	0.49700	0.49423
6	sibsp (+)	667.234	705.794	-113874510.309	0.50175	0.49846
7	deck_D (+)	663.423	706.804	-115313992.417	0.50490	0.50108
8	embarked_Q (+)	661.553	709.753	-116277909.273	0.50699	0.50264
9	embarked_Q (-)	663.423	706.804	-115313992.417	0.50490	0.50108

Final Model Output

Model Summary

R	0.711	RMSE	0.344
R-Squared	0.505	MSE	0.118
Adj. R-Squared	0.501	Coef. Var	87.229
Pred R-Squared	0.495	AIC	663.423
MAE	0.261	SBC	706.804

RMSE: Root Mean Square Error

MSE: Mean Square Error

MAE: Mean Absolute Error

AIC: Akaike Information Criteria

SBC: Schwarz Bayesian Criteria

ANOVA

	Sum of Squares	DF	Mean Square	F	Sig.
Regression	110.647	7	15.807	132.282	0.0000
Residual	108.500	908	0.119		
Total	219.147	915			

Parameter Estimates

model	Beta	Std. Error	Std. Beta	t	Sig	lower	upper
(Intercept)	1.128	0.044		25.918	0.000	1.043	1.213
sex_M	-0.465	0.024	-0.458	-19.077	0.000	-0.513	-0.418
deck_F	-0.450	0.031	-0.450	-14.496	0.000	-0.511	-0.389
age	-0.006	0.001	-0.170	-6.842	0.000	-0.008	-0.005
pclass_2	0.156	0.028	0.132	5.491	0.000	0.100	0.211

deck_E	0.165	0.043	0.105	3.842	0.000	0.081	0.249
sibsp	-0.032	0.011	-0.071	-2.957	0.003	-0.054	-0.011
deck_D	-0.116	0.048	-0.064	-2.404	0.016	-0.211	-0.021

```
# Fit model with stepwise parms only:
```

```
binary.model.stepwise = glm(survived ~ sex_M + deck_F + age + deck_E +
                             pclass_2 + pclass_1 + sibsp + fare,
                             family = binomial, data = train)
summary(binary.model.stepwise)
```

Call:

```
glm(formula = survived ~ sex_M + deck_F + age + deck_E + pclass_2 +
     pclass_1 + sibsp + fare, family = binomial, data = train)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	4.027276	0.443292	9.085	< 2e-16 ***
sex_M	-3.228025	0.242958	-13.286	< 2e-16 ***
deck_F	-2.649341	0.335751	-7.891	3.00e-15 ***
age	-0.054754	0.008774	-6.241	4.36e-10 ***
deck_E	1.823622	0.376762	4.840	1.30e-06 ***
pclass_2	1.472120	0.270394	5.444	5.20e-08 ***
pclass_1	0.531123	0.376285	1.411	0.15810
sibsp	-0.324078	0.115699	-2.801	0.00509 **
fare	0.001937	0.002089	0.927	0.35379

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 1230.15 on 915 degrees of freedom
 Residual deviance: 652.91 on 907 degrees of freedom
 AIC: 670.91

Number of Fisher Scoring iterations: 6

```
# Removing fare from stepwise model since its insig.:
```

```
binary.model.stepwise = update(binary.model.stepwise, paste(". ~ . -", 'fare'))
summary(binary.model.stepwise)
```

Call:

```
glm(formula = survived ~ sex_M + deck_F + age + deck_E + pclass_2 +
     pclass_1 + sibsp, family = binomial, data = train)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	4.075133	0.441121	9.238	< 2e-16 ***
sex_M	-3.250356	0.241996	-13.431	< 2e-16 ***

```

deck_F      -2.671528    0.335679   -7.959 1.74e-15 ***
age         -0.054914    0.008763   -6.266 3.70e-10 ***
deck_E       1.802546    0.376305    4.790 1.67e-06 ***
pclass_2     1.497212    0.269700    5.551 2.83e-08 ***
pclass_1     0.658065    0.351293    1.873 0.06103 .
sibsp       -0.305711    0.113440   -2.695 0.00704 **
---

```

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

```

Null deviance: 1230.15 on 915 degrees of freedom
Residual deviance: 653.81 on 908 degrees of freedom
AIC: 669.81

```

Number of Fisher Scoring iterations: 6

```

#####
#                               Model with influential points removed                               #
#####
# Applying function to remove influential points via Cooks Distance:
filtering.result = remove_cooks_outliers(survived ~ ., data = train)
cat(nrow(filtering.result$high_cd_data), 'rows were identified as outliers')

```

0 rows were identified as outliers

```

# plot(logmod)

# Since there were now rows identified as outliers, then this model will be the
# same as the initial binary model and need no be considered in the final
# model compare.

#####
#                               Model with high VIF preds removed                               #
#####
library(car)
vif(lmod)

```

	age	sibsp	parch	fare	pclass_1	pclass_2	sex_M
	1.493399	1.258340	1.281859	1.727034	4.267734	1.560100	1.604117
embarked_C	embarked_Q	deck_A	deck_B	deck_C	deck_D	deck_E	
	1.513511	1.397623	4.532575	4.848833	7.604802	6.662169	6.091048
deck_F							
	14.919819						

```

vif.model = glm(survived ~ . -deck_F, family = binomial, data = train)
vif(vif.model)

```

	age	sibsp	parch	fare	pclass_1	pclass_2	sex_M
	1.481259	1.236486	1.275733	1.729266	4.369376	1.472982	1.452735
embarked_C	embarked_Q	deck_A	deck_B	deck_C	deck_D	deck_E	
	1.511473	1.348123	1.988884	2.269651	2.863440	1.737448	1.558559

```
# Removing deck_F from the model eliminates the multicollinearity completely.
```

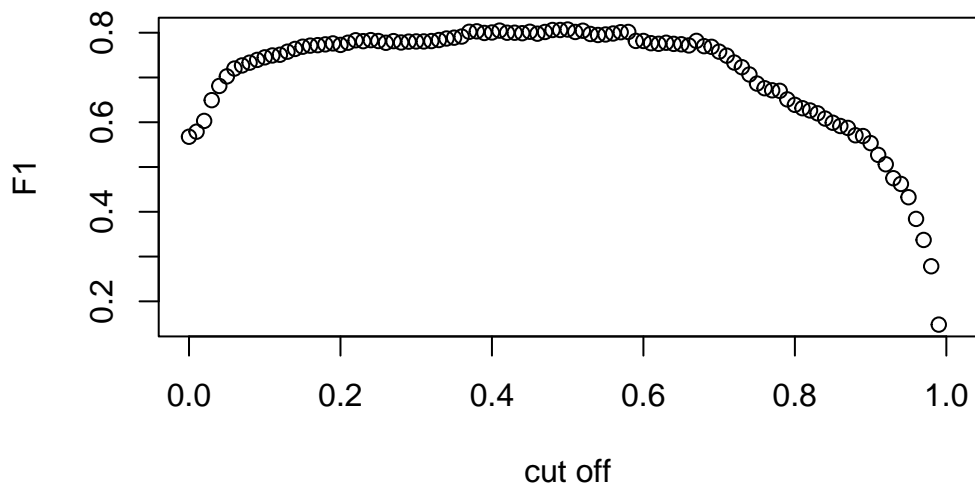
```
#####  
#                               Function to calc. cutoff                               #  
#####  
cutoff.prg<-function(pred,act){  
# pred<-predicted_probabilities  
# act<-true_labels  
p<-seq(0,1,0.01)  
n<-length(p)  
out<-matrix(0,nrow=n,ncol=12)  
for(i in 1:n){  
predictions <- ifelse(pred >p[i], 1, 0)  
confusion_matrix <- confusionMatrix(as.factor(predictions),as.factor(act),mode="prec_recall", position="right")  
out[i,]<-cbind(p=p[i],t(confusion_matrix[[4]]))  
}  
dimnames(out)[[2]]<-c("p","Sensitivity","Specificity","Pos Pred Value","Neg Pred Value","Precision","Recall")  
out  
}
```

```
# Finding the optimal cutoff  
observations = train$survived  
prob <- predict(lmod, train, type="response")  
  
test_cutoff<-cutoff.prg(prob,observations)
```

```
Warning in confusionMatrix.default(as.factor(predictions), as.factor(act), :  
Levels are not in the same order for reference and data. Refactoring data to  
match.
```

```
Warning in confusionMatrix.default(as.factor(predictions), as.factor(act), :  
Levels are not in the same order for reference and data. Refactoring data to  
match.
```

```
plot(test_cutoff[,1],test_cutoff[,8],xlab="cut off",ylab="F1")
```



```
optimal.cutoff = test_cutoff[which.max(test_cutoff[,8]),]
optimal.cutoff[1]
```

```
p
0.5
```

```
observations.train = train$survived
observations.test = test$survived

# Confusion matrix on base log model on train data
y_hat_prob = predict(lmod, train, type="response")
predictions.binary.model.train <- ifelse(y_hat_prob > optimal.cutoff[1], 1, 0)
confusion.matrix.binary.model.train <- confusionMatrix(as.factor(predictions.binary.model.train), as.factor(observations.train))
base.model.accuracy = confusion.matrix.binary.model.train$overall['Accuracy']
base.model.f1 = confusion.matrix.binary.model.train$byClass['F1']
base.model.train.summary = data.frame(
  Accuracy = base.model.accuracy,
  F1 = base.model.f1
)
row.names(base.model.train.summary) <- 'base.model.train'

# Confusion matrix on base log model on test data
y_hat_prob = predict(lmod, test, type="response")

predictions.binary.model.test <- ifelse(y_hat_prob > optimal.cutoff[1], 1, 0)

confusion.matrix.binary.model.test <- confusionMatrix(as.factor(predictions.binary.model.test), as.factor(observations.test))

base.model.accuracy.test = confusion.matrix.binary.model.test$overall['Accuracy']
base.model.f1.test = confusion.matrix.binary.model.test$byClass['F1']

base.model.test.summary = data.frame(
```



```

    Accuracy = base.model.accuracy.test,
    F1 = base.model.f1.test
)
row.names(base.model.test.summary) <- 'base.model.test'

# Confusion matrix on stepwise log model on train data
y_hat_prob = predict(binary.model.stepwise, train, type="response")
predictions.binary.model.step.train <- ifelse(y_hat_prob > optimal.cutoff[1], 1, 0)
confusion.matrix.binary.model.step.train <- confusionMatrix(as.factor(predictions.binary.model.step.train), train$y)
stepwise.model.accuracy = confusion.matrix.binary.model.step.train$overall['Accuracy']
stepwise.model.f1 = confusion.matrix.binary.model.step.train$byClass['F1']
stepwise.model.train.summary = data.frame(
  Accuracy = stepwise.model.accuracy,
  F1 = stepwise.model.f1
)
row.names(stepwise.model.train.summary) <- 'stepwise.model.train'

# Confusion matrix on stepwise log model on test data
y_hat_prob = predict(binary.model.stepwise, test, type="response")
predictions.binary.model.step.test <- ifelse(y_hat_prob > optimal.cutoff[1], 1, 0)
confusion.matrix.binary.model.step.test <- confusionMatrix(as.factor(predictions.binary.model.step.test), test$y)
stepwise.model.accuracy = confusion.matrix.binary.model.step.test$overall['Accuracy']
stepwise.model.f1 = confusion.matrix.binary.model.step.test$byClass['F1']
stepwise.model.test.summary = data.frame(
  Accuracy = stepwise.model.accuracy,
  F1 = stepwise.model.f1
)
row.names(stepwise.model.test.summary) <- 'stepwise.model.test'

# Confusion matrix on log model w/ insig. pred. removed
y_hat_prob = predict(binary.model.filtered, train, type="response")
predictions.binary.model.filtered.train <- ifelse(y_hat_prob > optimal.cutoff[1], 1, 0)
confusion.matrix.binary.model.filtered.train <- confusionMatrix(as.factor(predictions.binary.model.filtered.train), train$y)
filtered.model.accuracy = confusion.matrix.binary.model.filtered.train$overall['Accuracy']
filtered.model.f1 = confusion.matrix.binary.model.filtered.train$byClass['F1']
filtered.model.train.summary = data.frame(
  Accuracy = filtered.model.accuracy,
  F1 = filtered.model.f1
)
row.names(filtered.model.train.summary) <- 'filtered.model.train'

# Confusion matrix on log model w/ insig. pred. removed
y_hat_prob = predict(binary.model.filtered, test, type="response")

predictions.binary.model.filtered.test <- ifelse(y_hat_prob > optimal.cutoff[1], 1, 0)

confusion.matrix.binary.model.filtered.test <- confusionMatrix(as.factor(predictions.binary.model.filtered.test), test$y)

filtered.model.accuracy = confusion.matrix.binary.model.filtered.test$overall['Accuracy']
filtered.model.f1 = confusion.matrix.binary.model.filtered.test$byClass['F1']

filtered.model.test.summary = data.frame(

```

```

    Accuracy = filtered.model.accuracy,
    F1 = filtered.model.f1
)
row.names(filtered.model.test.summary) <- 'filtered.model.test'

# Confusion matrix on log model w/ high vif pred. removed
y_hat_prob = predict(vif.model, train, type="response")
predictions.binary.model.vif.train <- ifelse(y_hat_prob > optimal.cutoff[1], 1, 0)
confusion.matrix.binary.model.vif.train <- confusionMatrix(as.factor(predictions.binary.model.vif.train),
vif.model.accuracy = confusion.matrix.binary.model.vif.train$overall['Accuracy']
vif.model.f1 = confusion.matrix.binary.model.vif.train$byClass['F1']
vif.model.train.summary = data.frame(
    Accuracy = vif.model.accuracy,
    F1 = vif.model.f1
)
row.names(vif.model.train.summary) <- 'vif.model.train'

# Confusion matrix on log model w/ high vif pred. removed
y_hat_prob = predict(vif.model, test, type="response")
predictions.binary.model.vif.test <- ifelse(y_hat_prob > optimal.cutoff[1], 1, 0)
confusion.matrix.binary.model.vif.test <- confusionMatrix(as.factor(predictions.binary.model.vif.test),
vif.model.accuracy = confusion.matrix.binary.model.vif.test$overall['Accuracy']
vif.model.f1 = confusion.matrix.binary.model.vif.test$byClass['F1']
vif.model.test.summary = data.frame(
    Accuracy = vif.model.accuracy,
    F1 = vif.model.f1
)
row.names(vif.model.test.summary) <- 'vif.model.test'

data.frame(rbind(base.model.train.summary, base.model.test.summary,
    stepwise.model.train.summary, stepwise.model.test.summary,
    filtered.model.train.summary, filtered.model.test.summary,
    vif.model.train.summary, vif.model.test.summary))

```

Table 3: Comparison of all log models performance

	Accuracy	F1
base.model.train	0.8471616	0.8071625
base.model.test	0.8367347	0.7730496
stepwise.model.train	0.8220524	0.7687943
stepwise.model.test	0.8418367	0.7669173
filtered.model.train	0.8471616	0.8066298
filtered.model.test	0.8367347	0.7730496
vif.model.train	0.8329694	0.7877947
vif.model.test	0.8367347	0.7746479

7 Challenger Models

Build an alternative model based on one of the following approaches to predict survival as applicable: logistic regression, decision tree, NN, or SVM, Poisson regression or negative binomial. Check the applicable model

assumptions. Apply in-sample and out-of-sample testing, back testing and review the comparative goodness of fit of the candidate models. Describe step by step your procedure to get to the best model and why you believe it is fit for purpose.

```
library(MASS)

# Train base Poisson model:
poissonReg <- glm(survived ~ .,family=poisson, train)
summary(poissonReg)
```

Call:

```
glm(formula = survived ~ ., family = poisson, data = train)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)	
(Intercept)	7.519e-01	3.472e-01	2.166	0.030341	*
age	-1.668e-02	4.337e-03	-3.847	0.000119	***
sibsp	-7.014e-02	6.994e-02	-1.003	0.315889	
parch	1.132e-02	6.704e-02	0.169	0.865869	
fare	9.688e-05	1.025e-03	0.095	0.924696	
pclass_1	3.485e-01	1.987e-01	1.754	0.079421	.
pclass_2	5.468e-01	1.515e-01	3.609	0.000307	***
sex_M	-1.088e+00	1.199e-01	-9.069	< 2e-16	***
embarked_C	7.474e-02	1.362e-01	0.549	0.583046	
embarked_Q	5.098e-01	2.183e-01	2.335	0.019548	*
deck_A	-5.113e-01	4.521e-01	-1.131	0.258009	
deck_B	-5.677e-01	4.064e-01	-1.397	0.162402	
deck_C	-6.456e-01	3.912e-01	-1.650	0.098877	.
deck_D	-6.580e-01	3.820e-01	-1.722	0.084985	.
deck_E	-1.870e-01	3.526e-01	-0.530	0.595781	
deck_F	-1.540e+00	3.413e-01	-4.513	6.38e-06	***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for poisson family taken to be 1)

Null deviance: 672.00 on 915 degrees of freedom
Residual deviance: 400.84 on 900 degrees of freedom
AIC: 1158.8

Number of Fisher Scoring iterations: 5

```
# compare the fitted model to the null model
1-pchisq(672.34-390.72, length(coef(poissonReg)))
```

[1] 0

```
anova(poissonReg,test="Chi")
```

	Df	Deviance	Resid. Df	Resid. Dev	Pr(>Chi)
NULL	NA	NA	915	671.9954	NA
age	1	1.1812909	914	670.8141	0.2770933
sibsp	1	0.7166530	913	670.0975	0.3972440
parch	1	5.6694845	912	664.4280	0.0172625
fare	1	25.6933009	911	638.7347	0.0000004
pclass_1	1	27.9481171	910	610.7866	0.0000001
pclass_2	1	17.1639017	909	593.6227	0.0000343
sex_M	1	118.6444922	908	474.9782	0.0000000
embarked_C	1	1.6665466	907	473.3116	0.1967217
embarked_Q	1	0.0858234	906	473.2258	0.7695556
deck_A	1	0.0631207	905	473.1627	0.8016298
deck_B	1	0.0077426	904	473.1549	0.9298829
deck_C	1	0.0326392	903	473.1223	0.8566319
deck_D	1	0.5268888	902	472.5954	0.4679178
deck_E	1	57.6272332	901	414.9682	0.0000000
deck_F	1	14.1318730	900	400.8363	0.0001704

```
drop1(poissonReg, test="Chi")
```

	Df	Deviance	AIC	LRT	Pr(>Chi)
	NA	400.8363	1158.836	NA	NA
age	1	416.0444	1172.044	15.2080992	0.0000963
sibsp	1	401.8958	1157.896	1.0594618	0.3033380
parch	1	400.8646	1156.865	0.0282941	0.8664193
fare	1	400.8452	1156.845	0.0088926	0.9248705
pclass_1	1	403.9035	1159.903	3.0671630	0.0798888
pclass_2	1	413.5888	1169.589	12.7524939	0.0003555
sex_M	1	489.9790	1245.979	89.1427013	0.0000000
embarked_C	1	401.1362	1157.136	0.2998976	0.5839466
embarked_Q	1	405.8500	1161.850	5.0137077	0.0251474
deck_A	1	402.0796	1158.080	1.2432486	0.2648459
deck_B	1	402.6560	1158.656	1.8197039	0.1773485
deck_C	1	403.3130	1159.313	2.4766883	0.1155453
deck_D	1	403.5115	1159.511	2.6751598	0.1019251
deck_E	1	401.1067	1157.107	0.2703827	0.6030752
deck_F	1	414.9682	1170.968	14.1318730	0.0001704

```
# Testing against test data
predicted_probs <- predict(poissonReg, newdata = test, type = "response")

# Get the results of different cutoff values
testResult<-cutoff.prg(predicted_probs, test$survived)
```

Warning in confusionMatrix.default(as.factor(predictions), as.factor(act), :
Levels are not in the same order for reference and data. Refactoring data to
match.

Warning in confusionMatrix.default(as.factor(predictions), as.factor(act), :
Levels are not in the same order for reference and data. Refactoring data to

```

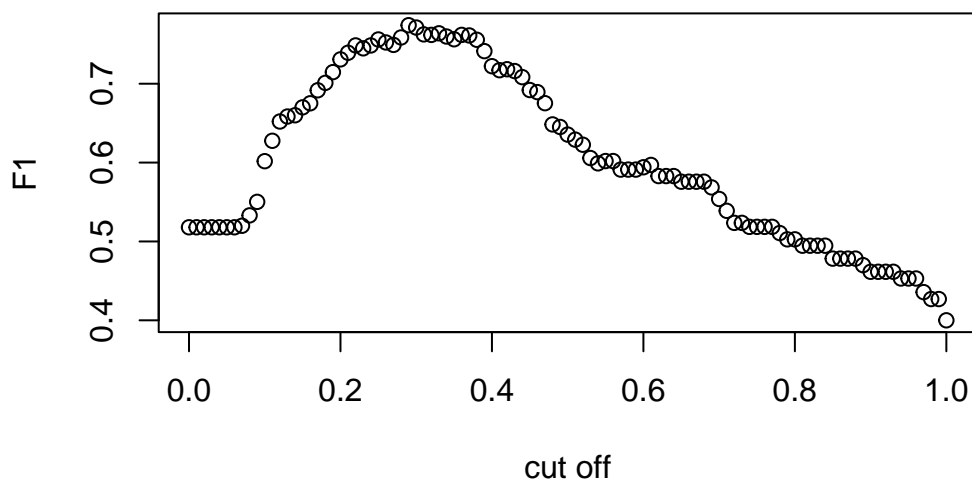
match.
Warning in confusionMatrix.default(as.factor(predictions), as.factor(act), :
Levels are not in the same order for reference and data. Refactoring data to
match.
Warning in confusionMatrix.default(as.factor(predictions), as.factor(act), :
Levels are not in the same order for reference and data. Refactoring data to
match.
Warning in confusionMatrix.default(as.factor(predictions), as.factor(act), :
Levels are not in the same order for reference and data. Refactoring data to
match.
Warning in confusionMatrix.default(as.factor(predictions), as.factor(act), :
Levels are not in the same order for reference and data. Refactoring data to
match.
Warning in confusionMatrix.default(as.factor(predictions), as.factor(act), :
Levels are not in the same order for reference and data. Refactoring data to
match.

```

```

# We get get the optimal p cut off value based on the maximum F1 score
plot(testResult[,1],testResult[,8],xlab="cut off",ylab="F1")

```



```

testResult[which.max(testResult[,8]),]

```

p	Sensitivity	Specificity
0.2900000	0.8759124	0.7921569
Pos Pred Value	Neg Pred Value	Precision
0.6936416	0.9223744	0.6936416
Recall	F1	Prevalence
0.8759124	0.7741935	0.3494898
Detection Rate	Detection Prevalence	Balanced Accuracy
0.3061224	0.4413265	0.8340346

8 Model Limitation and Assumptions

The champion model selected is the filtered logistic regression model due to its superior train accuracy (85.04%) and F1 score (0.8099). The filtered logistic regression model was developed through backward elimination procedure, applied to the full base model. This process iteratively removed predictors with high p-values until only statistically significant predictors remained, based on threshold ($p < 0.05$). The base model is used as the benchmark due to its similar performance and broader feature set.

“Survived = 5.56 - 0.055age - 0.23sibsp + 1pclass_1 + 1.9pclass_2 - 3.2sex_M + 0.84embarked_C + 0.86embarked_Q - 2.5deck_A - 1.9deck_B - 2.1deck_C - 3.1deck_D - 4.6deck_F”

To further evaluate and benchmark the logistic regression models, we computed RMSE, R^2 , and MAE using the predicted probabilities against the actual binary outcomes. While traditional R^2 is not appropriate for logistic models, these metrics provide insight into how well the predicted probabilities align with observed outcomes. The Filtered model appears to have higher R^2 ($0.4425 > 0.4342$), lower RMSE ($0.3945 < 0.3977$) and MAE ($0.1556 < 0.1582$). Additionally, model validation was done through train/test split with consistent performance across sets. The performance on test data and train data show minimal difference, indicating the model is robust, no overfitting problem.

To evaluate the stability and fit of the base and filtered logistic regression models, we conducted residual-based tests on their linear analogs. While logistic regression does not formally require normally distributed or homoscedastic residuals, these tests offer insights into model specification quality and residual behavior.

Breusch-Pagan Test: Both models exhibit statistically significant heteroskedasticity ($p < 0.01$)

Shapiro-Wilk Test: Both models show highly significant deviations from normality ($p < 0.01$)

Both models show similar residual behavior, with slight heteroskedasticity and non-normality. These findings do not violate logistic regression assumptions but suggest the model could be further improved through additional transformations or interaction terms.

Multicollinearity was significantly lower in the filtered model (max VIF: 3.69 vs. 16.63), supporting better generalization and interpretability. The VIF is less than 5, confirmed no severe multicollinearity.

```
library(car)

# Normality Test
# Breusch Pagan Test on Logistic Model
BPtest_basemodel <- ols_test_breusch_pagan(lmod)
BPtest_filtered.model <- ols_test_breusch_pagan(binary.model.filtered)

# Shapiro-Wilk test
SWtest_basemodel <- shapiro.test(residuals(lmod))
SWtest_filteredmodel <- shapiro.test(residuals(binary.model.filtered))

# Combine into unified rows
diagnostics_summary <- data.frame(
  Model = c("Base.model", "Filtered.model"),
  BP_Statistic = c(BPtest_basemodel$bp, BPtest_filtered.model$bp),
  BP_pvalue = c(BPtest_basemodel$p, BPtest_filtered.model$p),
  Shapiro_W = c(SWtest_basemodel$statistic, SWtest_filteredmodel$statistic),
  Shapiro_pvalue = c(SWtest_basemodel$p.value, SWtest_filteredmodel$p.value)
)
```

Table 6: Summary of Diagnostic tests on Filtered vs Base models

Model	BP_Statistic	BP_pvalue	Shapiro_W	Shapiro_pvalue
Base.model	10.11660	0.0014694	0.9752266	0
Filtered.model	10.31519	0.0013194	0.9751365	0

```
# Run VIF on the logistic model
ols_vif_tol(lmod)
```

Table 7: VIF Results for Base Logistic Regression Model

Variables	Tolerance	VIF
age	0.8112422	1.232677
sibsp	0.8133135	1.229538
parch	0.7793306	1.283152
fare	0.5101132	1.960349
pclass_1	0.2468157	4.051607
pclass_2	0.7943336	1.258917
sex_M	0.8512441	1.174751
embarked_C	0.7176662	1.393405
embarked_Q	0.8520166	1.173686
deck_A	0.2572099	3.887875
deck_B	0.1656314	6.037501
deck_C	0.1183078	8.452526
deck_D	0.1511669	6.615204
deck_E	0.1232885	8.111060
deck_F	0.0562819	17.767709

```
ols_vif_tol(binary.model.filtered)
```

Table 8: VIF Results for Filtered Logistic Regression Model

Variables	Tolerance	VIF
age	0.8251918	1.211839
sibsp	0.9323496	1.072559
pclass_1	0.2772087	3.607391
pclass_2	0.8044472	1.243090
sex_M	0.9085314	1.100677
embarked_C	0.7371420	1.356591
embarked_Q	0.8625868	1.159304
deck_A	0.6525446	1.532462
deck_B	0.5110353	1.956812
deck_C	0.4334379	2.307135
deck_D	0.5807436	1.721930
deck_F	0.3568101	2.802611

```

# Base Model - Train
observations.train <- train$survived
y_hat_base_train <- predict(lmod, train, type = "response")
pred_base_train <- ifelse(y_hat_base_train > optimal.cutoff[1], 1, 0)
ModelTrain_base <- data.frame(obs = observations.train, pred = pred_base_train)
log.train.base <- defaultSummary(ModelTrain_base)

# Base Model - Test
observations.test <- test$survived
y_hat_base_test <- predict(lmod, test, type = "response")
pred_base_test <- ifelse(y_hat_base_test > optimal.cutoff[1], 1, 0)
ModelTest_base <- data.frame(obs = observations.test, pred = pred_base_test)
log.test.base <- defaultSummary(ModelTest_base)

# Base Model - Train
observations.train <- train$survived
y_hat_base_train <- predict(binary.model.filtered, train, type = "response")
pred_base_train <- ifelse(y_hat_base_train > optimal.cutoff[1], 1, 0)
ModelTrain_base <- data.frame(obs = observations.train, pred = pred_base_train)
filtered.log.train.base <- defaultSummary(ModelTrain_base)

# Base Model - Test
observations.test <- test$survived
y_hat_base_test <- predict(binary.model.filtered, test, type = "response")
pred_base_test <- ifelse(y_hat_base_test > optimal.cutoff[1], 1, 0)
ModelTest_base <- data.frame(obs = observations.test, pred = pred_base_test)
filtered.test.base <- defaultSummary(ModelTest_base)

data.frame(rbind(log.train.base, log.test.base, filtered.log.train.base, filtered.test.base))

```

Table 9: Comparison of stats between base and filtered log models:

	RMSE	Rsquared	MAE
log.train.base	0.3909456	0.4631895	0.1528384
log.test.base	0.4040610	0.4177733	0.1632653
filtered.log.train.base	0.3909456	0.4627873	0.1528384
filtered.test.base	0.4040610	0.4177733	0.1632653

9 Ongoing Model Monitoring Plan

In order to maintain the effectiveness of the model, we would need to continue to test it on new data. Since the Titanic was a rare event, we do not have a lot of new data to test on the model, but we can still be prepared in case new data were to become available. The first step in monitoring the model is to determine specific thresholds that we expect the model to stay above. We would want the model to maintain certain R^2 , RMSE, and MAE values in order to determine that the model is working correctly. One of the biggest concerns with our model is data drift. Since the Titanic sank over 100 years ago, the data that we are using from the model may not align with today relevant to ship travel today.

10 Conclusion

Summarize your results here. What is the best model for the data and why?

Appendix A: Check if 'sibsp' and 'parch' should be continuous or categorical

We don't see significant improvement between modeling these predictors as continuous or categorical, therefore we decided to leave them as continuous.

```
library(car)
data.clean.ap1 = odata[, !(names(odata) %in% c("name", "ticket", "boat", "body", "home.dest"))]

#####
#                               Data Augmentation                               #
#####
#Extract deck letter from cabin
data.clean.ap1$deck <- substr(data.clean.ap1$cabin, 1,1)

# Remove cabin col:
data.clean.ap1$cabin <- NULL

#####
#                               Imputing data                               #
#####

# ---- Age----
#Replace NAs in age column with Median value
median_age <- median(data.clean.ap1$age, na.rm = TRUE)
data.clean.ap1 <- data.clean.ap1 %>%
  mutate(age = ifelse(is.na(age), median_age, age))

# ---- deck----
# For deck, since its a category, we decided to use KNN to impute the column:

# Install if not already installed
# install.packages("VIM")
library(VIM)

# Replace "" with NA in the 'deck' column
data.clean.ap1$deck[data.clean.ap1$deck == ""] <- NA

# Convert 'cabin' to factor
data.clean.ap1$deck <- as.factor(data.clean.ap1$deck)

# Apply kNN imputation just to Cabin column
data.clean.ap1 <- kNN(data.clean.ap1, variable = "deck", k = 5)

# Check that NAs were imputed
# sum(is.na(data.clean$deck))          # Original
# sum(is.na(data.clean.imputed$deck)) # After
```

```

# Remove indicator col:
data.clean.ap1$deck_imp <- NULL

#####
#                               Dummify Cat. cols                               #
#####

# Dummifying pclass:
data.clean.ap1$pclass_1 = ifelse(data.clean.ap1$pclass == 1, 1, 0)
data.clean.ap1$pclass_2 = ifelse(data.clean.ap1$pclass == 2, 1, 0)

# Dummifying sex:
data.clean.ap1$sex_M = ifelse(data.clean.ap1$sex == 'male', 1, 0)

# Dummifying embarked:
data.clean.ap1$embarked_C = ifelse(data.clean.ap1$embarked == 'C', 1, 0)
data.clean.ap1$embarked_Q = ifelse(data.clean.ap1$embarked == 'Q', 1, 0)

# Dummifying deck:
data.clean.ap1$deck_A = ifelse(data.clean.ap1$deck == 'A', 1, 0)
data.clean.ap1$deck_B = ifelse(data.clean.ap1$deck == 'B', 1, 0)
data.clean.ap1$deck_C = ifelse(data.clean.ap1$deck == 'C', 1, 0)
data.clean.ap1$deck_D = ifelse(data.clean.ap1$deck == 'D', 1, 0)
data.clean.ap1$deck_E = ifelse(data.clean.ap1$deck == 'E', 1, 0)
data.clean.ap1$deck_F = ifelse(data.clean.ap1$deck == 'F', 1, 0)
#data.clean.ap1$deck_G = ifelse(data.clean.ap1$deck == 'G', 1, 0) # removed due to causing issues

# Dummifying sibsp:
data.clean.ap1$sibsp_1 = ifelse(data.clean.ap1$sibsp == 1, 1, 0)
data.clean.ap1$sibsp_2 = ifelse(data.clean.ap1$sibsp == 2, 1, 0)
data.clean.ap1$sibsp_3 = ifelse(data.clean.ap1$sibsp == 3, 1, 0)
data.clean.ap1$sibsp_4 = ifelse(data.clean.ap1$sibsp == 4, 1, 0)
data.clean.ap1$sibsp_5 = ifelse(data.clean.ap1$sibsp == 5, 1, 0)
#data.clean.ap1$sibsp_8 = ifelse(data.clean.ap1$sibsp == 8, 1, 0) # removed due to causing issues

# Dummifying parch:
data.clean.ap1$parch_1 = ifelse(data.clean.ap1$parch == 1, 1, 0)
data.clean.ap1$parch_2 = ifelse(data.clean.ap1$parch == 2, 1, 0)
data.clean.ap1$parch_3 = ifelse(data.clean.ap1$parch == 3, 1, 0)
data.clean.ap1$parch_4 = ifelse(data.clean.ap1$parch == 4, 1, 0)
data.clean.ap1$parch_5 = ifelse(data.clean.ap1$parch == 5, 1, 0)
data.clean.ap1$parch_6 = ifelse(data.clean.ap1$parch == 6, 1, 0)
#data.clean.ap1$parch_9 = ifelse(data.clean.ap1$parch == 9, 1, 0) # removed due to causing issues

# Removing Dummified cols:
data.clean.ap1 = subset(data.clean.ap1, select = -c(pclass, sex, embarked, deck))#, sibsp, parch))

data.clean.ap1 = na.omit(data.clean.ap1)

cat(nrow(odata) - nrow(data.clean.ap1), 'rows were removed from original dataset')

```

2 rows were removed from original dataset

```

set.seed(567)
train_indices_ap1 = sample(1 : nrow(data.clean.ap1), size = 0.7005*nrow(data.clean.ap1), replace = FALSE)
train.ap1 = data.clean.ap1[train_indices_ap1,]
test.ap1 = data.clean.ap1[-train_indices_ap1,]
cat("We are using:", nrow(train.ap1)/nrow(data.clean.ap1) * 100, '% of the data for training')

```

We are using: 70.03058 % of the data for training

```

mulvar_model.ap1 <- lm(survived ~ ., data = train.ap1)
summary(mulvar_model.ap1)

```

Call:

```
lm(formula = survived ~ ., data = train.ap1)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-1.2866	-0.1941	-0.0224	0.1911	0.9739

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	1.1796068	0.1203338	9.803	< 2e-16	***
age	-0.0064249	0.0010420	-6.166	1.06e-09	***
sibsp	-0.0401162	0.0221954	-1.807	0.07104	.
parch	-0.0274112	0.0270379	-1.014	0.31095	
fare	0.0004356	0.0002985	1.459	0.14481	
pclass_1	0.1097829	0.0542228	2.025	0.04320	*
pclass_2	0.1381701	0.0312063	4.428	1.07e-05	***
sex_M	-0.4129619	0.0252816	-16.335	< 2e-16	***
embarked_C	0.0461410	0.0323127	1.428	0.15366	
embarked_Q	0.1210184	0.0426024	2.841	0.00461	**
deck_A	-0.3494376	0.1454048	-2.403	0.01646	*
deck_B	-0.3123626	0.1372942	-2.275	0.02314	*
deck_C	-0.3040370	0.1331009	-2.284	0.02259	*
deck_D	-0.3432119	0.1290127	-2.660	0.00795	**
deck_E	0.0041655	0.1243264	0.034	0.97328	
deck_F	-0.5934388	0.1184958	-5.008	6.63e-07	***
sibsp_1	0.0567316	0.0355896	1.594	0.11128	
sibsp_2	0.1212082	0.0797596	1.520	0.12895	
sibsp_3	-0.1501858	0.1195550	-1.256	0.20937	
sibsp_4	-0.2298081	0.1296372	-1.773	0.07662	.
sibsp_5	-0.2293432	0.2019073	-1.136	0.25631	
parch_1	0.1426458	0.0454040	3.142	0.00174	**
parch_2	0.1456666	0.0715147	2.037	0.04196	*
parch_3	0.2704118	0.1739380	1.555	0.12039	
parch_4	-0.0799175	0.2038416	-0.392	0.69511	
parch_5	0.0679073	0.2385574	0.285	0.77597	
parch_6	0.0143690	0.2895211	0.050	0.96043	

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.34 on 889 degrees of freedom
Multiple R-squared: 0.5273, Adjusted R-squared: 0.5135
F-statistic: 38.14 on 26 and 889 DF, p-value: < 2.2e-16

```
vif(mulvar_model.ap1)
```

	age	sibsp	parch	fare	pclass_1	pclass_2	sex_M
	1.460949	3.467000	4.494487	1.922834	4.492433	1.297793	1.164100
	embarked_C	embarked_Q	deck_A	deck_B	deck_C	deck_D	deck_E
	1.437150	1.172137	5.135825	8.715727	13.411354	9.672668	12.015361
	deck_F	sibsp_1	sibsp_2	sibsp_3	sibsp_4	sibsp_5	parch_1
	26.716246	1.876291	1.441990	1.464254	1.863037	1.404398	1.911901
	parch_2	parch_3	parch_4	parch_5	parch_6		
	3.083293	1.301393	1.431435	1.472004	1.446997		

```
lmod.ap1 <- glm(as.factor(survived) ~ ., family = binomial, data = train.ap1)
summary(lmod.ap1)
```

Call:

```
glm(formula = as.factor(survived) ~ ., family = binomial, data = train.ap1)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)	
(Intercept)	5.253e+00	1.384e+00	3.797	0.000147	***
age	-5.481e-02	9.965e-03	-5.500	3.79e-08	***
sibsp	-1.946e+00	1.291e+02	-0.015	0.987976	
parch	-1.642e+00	1.623e+02	-0.010	0.991928	
fare	1.057e-03	2.428e-03	0.435	0.663393	
pclass_1	1.156e+00	4.819e-01	2.398	0.016480	*
pclass_2	1.496e+00	3.095e-01	4.834	1.34e-06	***
sex_M	-2.982e+00	2.464e-01	-12.103	< 2e-16	***
embarked_C	6.410e-01	2.868e-01	2.235	0.025424	*
embarked_Q	1.316e+00	3.735e-01	3.523	0.000427	***
deck_A	-2.866e+00	1.503e+00	-1.907	0.056492	.
deck_B	-2.498e+00	1.479e+00	-1.689	0.091220	.
deck_C	-2.624e+00	1.445e+00	-1.817	0.069262	.
deck_D	-2.971e+00	1.415e+00	-2.100	0.035711	*
deck_E	3.196e-01	1.404e+00	0.228	0.819965	
deck_F	-4.753e+00	1.358e+00	-3.499	0.000467	***
sibsp_1	2.065e+00	1.291e+02	0.016	0.987240	
sibsp_2	4.356e+00	2.583e+02	0.017	0.986546	
sibsp_3	3.607e+00	3.874e+02	0.009	0.992572	
sibsp_4	4.557e+00	5.166e+02	0.009	0.992962	
sibsp_5	-6.723e+00	1.240e+03	-0.005	0.995674	
parch_1	2.597e+00	1.623e+02	0.016	0.987237	
parch_2	4.192e+00	3.247e+02	0.013	0.989698	
parch_3	6.236e+00	4.870e+02	0.013	0.989785	
parch_4	6.227e+00	6.494e+02	0.010	0.992349	
parch_5	8.544e+00	8.117e+02	0.011	0.991602	
parch_6	-4.112e+00	1.770e+03	-0.002	0.998146	

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 1223.12 on 915 degrees of freedom
Residual deviance: 602.86 on 889 degrees of freedom
AIC: 656.86

Number of Fisher Scoring iterations: 15

```
vif(lmod.ap1)
```

age	sibsp	parch	fare	pclass_1	pclass_2
1.653195e+00	7.591027e+05	1.634339e+06	1.757536e+00	4.949336e+00	1.555093e+00
sex_M	embarked_C	embarked_Q	deck_A	deck_B	deck_C
1.404155e+00	1.549826e+00	1.415978e+00	1.012962e+01	1.475410e+01	2.517074e+01
deck_D	deck_E	deck_F	sibsp_1	sibsp_2	sibsp_3
1.980371e+01	1.464836e+01	4.346983e+01	3.294475e+05	1.808931e+05	1.937224e+05
sibsp_4	sibsp_5	parch_1	parch_2	parch_3	parch_4
1.974919e+05	1.372219e+00	3.139987e+05	6.986245e+05	2.469813e+05	2.280957e+05
parch_5	parch_6				
4.056249e+05	1.434520e+00				

```
y_hat_mulvar_train.ap1<-predict(mulvar_model.ap1, data = train.ap1)
predictions_train.ap1 <- ifelse(y_hat_mulvar_train.ap1 > 0.5, 1, 0)
ModelTrain_mulvar.ap1<-data.frame(obs = train.ap1$survived, pred=predictions_train.ap1)
linear.train.ap1 <- defaultSummary(ModelTrain_mulvar.ap1)

y_hat_mulvar_test.ap1<-predict(mulvar_model.ap1, newdata = test.ap1)
predictions_test.ap1 <- ifelse(y_hat_mulvar_test.ap1 > 0.5, 1, 0)
ModelTest_mulvar.ap1<-data.frame(obs = test.ap1$survived, pred=predictions_test.ap1)
linear.test.ap1 <- defaultSummary(ModelTest_mulvar.ap1)

y_hat_log_train.ap1<-predict(lmod.ap1, data = train.ap1)
predictions_log_train.ap1 <- ifelse(y_hat_log_train.ap1 > 0.5, 1, 0)
ModelTrain_lmod.ap1<-data.frame(obs = train.ap1$survived, pred=predictions_log_train.ap1)
log.train.ap1 <- defaultSummary(ModelTrain_lmod.ap1)

y_hat_log_test.ap1<-predict(lmod.ap1, newdata = test.ap1)
predictions_log_test.ap1 <- ifelse(y_hat_log_test.ap1 > 0.5, 1, 0)
ModelTest_lmod.ap1<-data.frame(obs = test.ap1$survived, pred=predictions_log_test.ap1)
log.test.ap1 <- defaultSummary(ModelTest_lmod.ap1)

data.frame(rbind(linear.train.ap1,linear.test.ap1,log.train.ap1,log.test.ap1))
```

	RMSE	Rsquared	MAE
linear.train.ap1	0.3796114	0.4824536	0.1441048
linear.test.ap1	0.3976975	0.4365256	0.1581633

	RMSE	Rsquared	MAE
log.train.ap1	0.3964912	0.4429840	0.1572052
log.test.ap1	0.4008919	0.4214242	0.1607143

```
confusion_matrix_mulvar_train.ap1 <- confusionMatrix(as.factor(predictions_train.ap1), as.factor(test.ap1))
confusion_matrix_mulvar_train.ap1
```

Confusion Matrix and Statistics

```

      Reference
Prediction  0   1
      0 502  73
      1  59 282

      Accuracy : 0.8559
      95% CI : (0.8315, 0.878)
No Information Rate : 0.6124
P-Value [Acc > NIR] : <2e-16

      Kappa : 0.6942

McNemar's Test P-Value : 0.2578

      Precision : 0.8270
      Recall : 0.7944
      F1 : 0.8103
      Prevalence : 0.3876
      Detection Rate : 0.3079
      Detection Prevalence : 0.3723
      Balanced Accuracy : 0.8446

      'Positive' Class : 1

```

```
confusion_matrix_mulvar_test.ap1 <- confusionMatrix(as.factor(predictions_test.ap1), as.factor(test.ap1))
confusion_matrix_mulvar_test.ap1
```

Confusion Matrix and Statistics

```

      Reference
Prediction  0   1
      0 216  31
      1  31 114

      Accuracy : 0.8418
      95% CI : (0.8019, 0.8765)
No Information Rate : 0.6301
P-Value [Acc > NIR] : <2e-16

      Kappa : 0.6607

```

Mcnemar's Test P-Value : 1

Precision : 0.7862
Recall : 0.7862
F1 : 0.7862
Prevalence : 0.3699
Detection Rate : 0.2908
Detection Prevalence : 0.3699
Balanced Accuracy : 0.8304

'Positive' Class : 1

```
confusion_matrix_log_train.ap1 <- confusionMatrix(as.factor(predictions_log_train.ap1), as.factor(tests_log_train.ap1))  
confusion_matrix_log_train.ap1
```

Confusion Matrix and Statistics

	Reference	
Prediction	0	1
0	523	106
1	38	249

Accuracy : 0.8428
95% CI : (0.8176, 0.8658)
No Information Rate : 0.6124
P-Value [Acc > NIR] : < 2.2e-16

Kappa : 0.6568

Mcnemar's Test P-Value : 2.36e-08

Precision : 0.8676
Recall : 0.7014
F1 : 0.7757
Prevalence : 0.3876
Detection Rate : 0.2718
Detection Prevalence : 0.3133
Balanced Accuracy : 0.8168

'Positive' Class : 1

```
confusion_matrix_log_test.ap1 <- confusionMatrix(as.factor(predictions_log_test.ap1), as.factor(tests_log_test.ap1))  
confusion_matrix_log_test.ap1
```

Confusion Matrix and Statistics

	Reference	
Prediction	0	1

0 226 42
1 21 103

Accuracy : 0.8393
95% CI : (0.7991, 0.8742)
No Information Rate : 0.6301
P-Value [Acc > NIR] : < 2e-16

Kappa : 0.6446

McNemar's Test P-Value : 0.01174

Precision : 0.8306
Recall : 0.7103
F1 : 0.7658
Prevalence : 0.3699
Detection Rate : 0.2628
Detection Prevalence : 0.3163
Balanced Accuracy : 0.8127

'Positive' Class : 1

```
library(car)
data.clean.ap2 = odata[, !(names(odata) %in% c("name", "ticket", "boat","body","home.dest"))]

#####
#                               Data Augmentation                               #
#####
#Extract deck letter from cabin
data.clean.ap2$deck <- substr(data.clean.ap2$cabin, 1,1)

# Remove cabin col:
data.clean.ap2$cabin <- NULL

#####
#                               Imputing data                               #
#####

# ---- Age----
#Replace NAs in age column with Median value
median_age <- median(data.clean.ap2$age, na.rm = TRUE)
data.clean.ap2 <- data.clean.ap2 %>%
  mutate(age = ifelse(is.na(age), median_age, age))

# ---- deck----
# For deck, since its a category, we decided to use KNN to impute the column:

# Install if not already installed
# install.packages("VIM")
library(VIM)

# Replace "" with NA in the 'deck' column
```



```

data.clean.ap2$deck[data.clean.ap2$deck == ""] <- NA

# Convert 'cabin' to factor
data.clean.ap2$deck <- as.factor(data.clean.ap2$deck)

# Apply kNN imputation just to Cabin column
data.clean.ap2 <- kNN(data.clean.ap2, variable = "deck", k = 5)

# Check that NAs were imputed
# sum(is.na(data.clean$deck))          # Original
# sum(is.na(data.clean.imputed$deck)) # After

# Remove indicator col:
data.clean.ap2$deck_imp <- NULL

#####
#                               Dummify Cat. cols                               #
#####

# Dummifying pclass:
data.clean.ap2$pclass_1 = ifelse(data.clean.ap2$pclass == 1, 1, 0)
data.clean.ap2$pclass_2 = ifelse(data.clean.ap2$pclass == 2, 1, 0)

# Dummifying sex:
data.clean.ap2$sex_M = ifelse(data.clean.ap2$sex == 'male', 1, 0)

# Dummifying embarked:
data.clean.ap2$embarked_C = ifelse(data.clean.ap2$embarked == 'C', 1, 0)
data.clean.ap2$embarked_Q = ifelse(data.clean.ap2$embarked == 'Q', 1, 0)

# Dummifying deck:
data.clean.ap2$deck_A = ifelse(data.clean.ap2$deck == 'A', 1, 0)
data.clean.ap2$deck_B = ifelse(data.clean.ap2$deck == 'B', 1, 0)
data.clean.ap2$deck_C = ifelse(data.clean.ap2$deck == 'C', 1, 0)
data.clean.ap2$deck_D = ifelse(data.clean.ap2$deck == 'D', 1, 0)
data.clean.ap2$deck_E = ifelse(data.clean.ap2$deck == 'E', 1, 0)
data.clean.ap2$deck_F = ifelse(data.clean.ap2$deck == 'F', 1, 0)
#data.clean.ap2$deck_G = ifelse(data.clean.ap2$deck == 'G', 1, 0) # removed due to causing issues

# Dummifying sibsp to 2 categories:
data.clean.ap2$sibsp_y = ifelse(data.clean.ap2$sibsp > 0, 1, 0)

# Dummifying parch to 2 categories:
data.clean.ap2$parch_y = ifelse(data.clean.ap2$parch > 0, 1, 0)

# Removing Dummified cols:
data.clean.ap2 = subset(data.clean.ap2, select = -c(pclass, sex, embarked, deck))#, sibsp, parch))

data.clean.ap2 = na.omit(data.clean.ap2)

cat(nrow(odata) - nrow(data.clean.ap2), 'rows were removed from original dataset')

```

2 rows were removed from original dataset

```
set.seed(567)
train_indices_ap2 = sample(1 : nrow(data.clean.ap2), size = 0.7005*nrow(data.clean.ap2), replace = FALSE)
train.ap2 = data.clean.ap2[train_indices_ap2,]
test.ap2 = data.clean.ap2[-train_indices_ap2,]
cat("We are using:", nrow(train.ap2)/nrow(data.clean.ap2) * 100, '% of the data for training')
```

We are using: 70.03058 % of the data for training

```
mulvar_model.ap2 <- lm(survived ~ ., data = train.ap2)
summary(mulvar_model.ap2)
```

Call:

```
lm(formula = survived ~ ., data = train.ap2)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-1.27243	-0.19451	-0.02769	0.19202	0.96436

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	1.1720948	0.1212992	9.663	< 2e-16 ***
age	-0.0061866	0.0010236	-6.044	2.20e-09 ***
sibsp	-0.0867184	0.0185029	-4.687	3.21e-06 ***
parch	-0.0487147	0.0214741	-2.269	0.023534 *
fare	0.0003848	0.0003005	1.281	0.200641
pclass_1	0.1055453	0.0557446	1.893	0.058629 .
pclass_2	0.1376903	0.0314471	4.378	1.34e-05 ***
sex_M	-0.4270550	0.0253448	-16.850	< 2e-16 ***
embarked_C	0.0585880	0.0319987	1.831	0.067439 .
embarked_Q	0.1061446	0.0425937	2.492	0.012881 *
deck_A	-0.3315651	0.1464353	-2.264	0.023797 *
deck_B	-0.2852592	0.1389479	-2.053	0.040362 *
deck_C	-0.2951467	0.1350558	-2.185	0.029120 *
deck_D	-0.3605291	0.1318846	-2.734	0.006386 **
deck_E	0.0078459	0.1256544	0.062	0.950226
deck_F	-0.5745493	0.1193029	-4.816	1.72e-06 ***
sibsp_y	0.1224343	0.0370630	3.303	0.000993 ***
parch_y	0.1573960	0.0479445	3.283	0.001067 **

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.3436 on 898 degrees of freedom

Multiple R-squared: 0.5124, Adjusted R-squared: 0.5031

F-statistic: 55.5 on 17 and 898 DF, p-value: < 2.2e-16

```
vif(mulvar_model.ap2)
```

age	sibsp	parch	fare	pclass_1	pclass_2	sex_M
1.380228	2.359204	2.776007	1.908310	4.649205	1.290432	1.145551
embarked_C	embarked_Q	deck_A	deck_B	deck_C	deck_D	deck_E
1.379990	1.147247	5.439805	8.740948	13.154666	9.148653	11.598396
deck_F	sibsp_y	parch_y				
26.237760	2.297316	3.213776				

```
lmod.ap2 <- glm(as.factor(survived) ~ ., family = binomial, data = train.ap2)
summary(lmod.ap2)
```

Call:

```
glm(formula = as.factor(survived) ~ ., family = binomial, data = train.ap2)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)	
(Intercept)	5.1275081	1.3016899	3.939	8.18e-05	***
age	-0.0558992	0.0094958	-5.887	3.94e-09	***
sibsp	-1.0266480	0.2640080	-3.889	0.000101	***
parch	-0.2791962	0.1939136	-1.440	0.149925	
fare	0.0007343	0.0024042	0.305	0.760042	
pclass_1	1.1421684	0.4862413	2.349	0.018825	*
pclass_2	1.4015232	0.2992294	4.684	2.82e-06	***
sex_M	-3.0275901	0.2424805	-12.486	< 2e-16	***
embarked_C	0.6732550	0.2814737	2.392	0.016762	*
embarked_Q	1.2284326	0.3679272	3.339	0.000841	***
deck_A	-2.5767916	1.4219365	-1.812	0.069960	.
deck_B	-2.2194317	1.4042525	-1.581	0.113991	
deck_C	-2.3552547	1.3682224	-1.721	0.085179	.
deck_D	-2.9570245	1.3451465	-2.198	0.027928	*
deck_E	0.4470949	1.3202938	0.339	0.734886	
deck_F	-4.4528239	1.2754924	-3.491	0.000481	***
sibsp_y	1.3882952	0.3947966	3.516	0.000437	***
parch_y	1.2304958	0.4199141	2.930	0.003386	**

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 1223.12 on 915 degrees of freedom
 Residual deviance: 620.68 on 898 degrees of freedom
 AIC: 656.68

Number of Fisher Scoring iterations: 6

```
vif(lmod.ap2)
```

age	sibsp	parch	fare	pclass_1	pclass_2	sex_M
1.619429	3.641267	2.989704	1.753639	5.122869	1.534060	1.410059
embarked_C	embarked_Q	deck_A	deck_B	deck_C	deck_D	deck_E
1.487757	1.363260	9.852461	13.562280	21.612617	16.823106	13.659175

deck_F	sibsp_y	parch_y
39.454519	3.482784	3.287406

```

y_hat_mulvar_train.ap2<-predict(mulvar_model.ap2, data = train.ap2)
predictions_train.ap2 <- ifelse(y_hat_mulvar_train.ap2 > 0.5, 1, 0)
ModelTrain_mulvar.ap2<-data.frame(obs = train.ap2$survived, pred=predictions_train.ap2)
linear.train.ap2 <- defaultSummary(ModelTrain_mulvar.ap2)

y_hat_mulvar_test.ap2<-predict(mulvar_model.ap2, newdata = test.ap2)
predictions_test.ap2 <- ifelse(y_hat_mulvar_test.ap2 > 0.5, 1, 0)
ModelTest_mulvar.ap2<-data.frame(obs = test.ap2$survived, pred=predictions_test.ap2)
linear.test.ap2 <- defaultSummary(ModelTest_mulvar.ap2)

y_hat_log_train.ap2<-predict(lmod.ap2, data = train.ap2)
predictions_log_train.ap2 <- ifelse(y_hat_log_train.ap2 > 0.5, 1, 0)
ModelTrain_lmod.ap2<-data.frame(obs = train.ap2$survived, pred=y_hat_log_train.ap2)
log.train.ap2 <- defaultSummary(ModelTrain_mulvar.ap2)

y_hat_log_test.ap2<-predict(lmod.ap2, newdata = test.ap2)
predictions_log_test.ap2 <- ifelse(y_hat_log_test.ap2 > 0.5, 1, 0)
ModelTest_lmod.ap2<-data.frame(obs = test.ap2$survived, pred=predictions_log_test.ap2)
log.test.ap2 <- defaultSummary(ModelTest_lmod.ap2)

data.frame(rbind(linear.train.ap2,linear.test.ap2,log.train.ap2,log.test.ap2))

```

	RMSE	Rsquared	MAE
linear.train.ap2	0.3909456	0.4570397	0.1528384
linear.test.ap2	0.4040610	0.4221825	0.1632653
log.train.ap2	0.3909456	0.4570397	0.1528384
log.test.ap2	0.4008919	0.4212961	0.1607143

```

confusion_matrix_mulvar_train.ap2 <- confusionMatrix(as.factor(predictions_train.ap2), as.factor(train.ap2$survived))
confusion_matrix_mulvar_train.ap2

```

Confusion Matrix and Statistics

	Reference	
Prediction	0	1
0	498	77
1	63	278

Accuracy : 0.8472
 95% CI : (0.8222, 0.8699)
 No Information Rate : 0.6124
 P-Value [Acc > NIR] : <2e-16

Kappa : 0.6757

McNemar's Test P-Value : 0.2719

Precision : 0.8152
Recall : 0.7831
F1 : 0.7989
Prevalence : 0.3876
Detection Rate : 0.3035
Detection Prevalence : 0.3723
Balanced Accuracy : 0.8354

'Positive' Class : 1

```
confusion_matrix_mulvar_test.ap2 <- confusionMatrix(as.factor(predictions_test.ap2), as.factor(test.ap2))  
confusion_matrix_mulvar_test.ap2
```

Confusion Matrix and Statistics

	Reference	
Prediction	0	1
0	215	32
1	32	113

Accuracy : 0.8367
95% CI : (0.7964, 0.8719)
No Information Rate : 0.6301
P-Value [Acc > NIR] : <2e-16

Kappa : 0.6498

McNemar's Test P-Value : 1

Precision : 0.7793
Recall : 0.7793
F1 : 0.7793
Prevalence : 0.3699
Detection Rate : 0.2883
Detection Prevalence : 0.3699
Balanced Accuracy : 0.8249

'Positive' Class : 1

```
confusion_matrix_log_train.ap2 <- confusionMatrix(as.factor(predictions_log_train.ap2), as.factor(test.ap2))  
confusion_matrix_log_train.ap2
```

Confusion Matrix and Statistics

	Reference	
Prediction	0	1
0	523	111

1 38 244

Accuracy : 0.8373
95% CI : (0.8118, 0.8607)
No Information Rate : 0.6124
P-Value [Acc > NIR] : < 2.2e-16

Kappa : 0.6439

McNemar's Test P-Value : 3.669e-09

Precision : 0.8652
Recall : 0.6873
F1 : 0.7661
Prevalence : 0.3876
Detection Rate : 0.2664
Detection Prevalence : 0.3079
Balanced Accuracy : 0.8098

'Positive' Class : 1

```
confusion_matrix_log_test.ap2 <- confusionMatrix(as.factor(predictions_log_test.ap2), as.factor(test_data_log_test.ap2))  
confusion_matrix_log_test.ap2
```

Confusion Matrix and Statistics

	Reference	
Prediction	0	1
0	227	43
1	20	102

Accuracy : 0.8393
95% CI : (0.7991, 0.8742)
No Information Rate : 0.6301
P-Value [Acc > NIR] : < 2.2e-16

Kappa : 0.6436

McNemar's Test P-Value : 0.005576

Precision : 0.8361
Recall : 0.7034
F1 : 0.7640
Prevalence : 0.3699
Detection Rate : 0.2602
Detection Prevalence : 0.3112
Balanced Accuracy : 0.8112

'Positive' Class : 1

References

National Oceanic and Atmospheric Administration (NOAA). (2023). *RMS titanic – history and significance*. <https://www.noaa.gov/office-of-general-counsel/gc-international-section/rms-titanic-history-and-significance>

Statology. (2025). *How to measure correlation between categorical variables*. <https://www.statology.org/correlation-between-categorical-variables/>